

Final Project Submission

Please fill out:

- Student name: MARTIN MWANGI
- Student pace: FULL TIME
- Scheduled project review date/time:
- Instructor name: ANTONNY MUIKO
- Blog post URL:

Business Problem

Voice of Kenya now sees all the big companies creating original video content and they want to get in on the fun. They have decided to create a new movie studio, but they don't know anything about creating movies. You are charged with exploring what types of films are currently doing the best at the box office. You must then translate those findings into actionable insights that the head of your company's new movie studio can use to help decide what type of films to create.

Business Understanding

The business problem presented by Voice of Kenya is to establish their own movie studio to compete within the movie market both locally and globally. For Voice of Kenya to succeed in this business, some questions need answers:

- Are movies really profitable, do they really make money
- What kind of the movie genre will be the most successful
- What is the budget of producing a movie from start to airing
- Who are my competitor and how established are they in the market

Answering these questions will help the Voice of Kenya to make decisions

Data Understanding

After carefully analysing the data provided in relation to the business problem and the business understanding question, I have selected the following datasets together with the columns that will be used in each particular dataset. The primary key is movie_title

1. [im.db.zip](#)(movie_basic i.e movie_title and genres)
2. [tn.movie_budgets.csv.gz](#)(movie_title, production_budget, domestic_gross, worldwide_gross)
3. [bom.movie_gross.csv.gz](#)(movie_title, studio)

Data preparation

The three datasets will go through a sanity check first that is data cleaning, it includes:

1. Converting some columns that are supposed to be numerical i.e production_budget, domestic_gross_y, worldwide_gross from object dtype to float dtype
2. Check the null or missing values and fill them, and drop where need be
3. Check and drop duplicates
4. Check and drop outliers

After the sanity check then i will merge the datasets 2 by 2 i.e will merge two datasets, then i will take the merged dataset and merge with the third dataset to have one dataset.

```
import pandas as pd
import sqlite3
import numpy as np
import seaborn as sns
import sklearn as sk
import statsmodels.api as sm
import scipy.stats as stats
import matplotlib.pyplot as plt
%matplotlib inline

import zipfile

# define the files path to extract and save the file after extract
zip_path = "zippedData/im.db.zip"
idmb_unzip = "zippedData/unzip idmb"
# unzip the datafile
with zipfile.ZipFile(zip_path, "r") as zip_ref:
    zip_ref.extractall(idmb_unzip)

# connecting our database
conn = sqlite3.connect("zippedData/unzip idmb/im.db")
# retriving our tables
movie_basics = pd.read_sql_query("SELECT* FROM movie_basics", conn)
movie_ratings = pd.read_sql_query("SELECT* FROM movie_ratings", conn)
# closing the connection

# connecting our datasets
# in the movie basic table i just need two columns the genre and
primary title
conn = sqlite3.connect("zippedData/unzip idmb/im.db")
q = """
SELECT
    primary_title,
    genres
FROM movie_basics;
"""
idmb_unzip = pd.read_sql(q, conn)

# cheking the info our movie basic column
idmb_unzip.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 146144 entries, 0 to 146143
Data columns (total 2 columns):
#   Column          Non-Null Count  Dtype
---  -
0   primary_title    146144 non-null  object
1   genres           140736 non-null  object
dtypes: object(2)
memory usage: 2.2+ MB

```

Data analysis

1. Checking missing values drop them or fill them
2. Checking for duplicates
3. Checking for outliers but in this dataset we have object dtype so we will omit checking the outliers

```

# checking of null values
idmb_unzip.isna().sum()

primary_title    0
genres          5408
dtype: int64

# our dataset contains 5408 missing values. This is a very small
percentage so we can drop the missing value rows
idmb_unzip1= idmb_unzip.dropna()
idmb_unzip1

```

```

                                primary_title
genres
0                                Sunghursh
Action,Crime,Drama
1          One Day Before the Rainy Season
Biography,Drama
2          The Other Side of the Wind
Drama
3          Sabse Bada Sukh
Comedy,Drama
4          The Wandering Soap Opera
Comedy,Drama,Fantasy
...
...
146138          The Secret of China
Adventure,History,War
146139          Kuambil Lagi Hatiku
Drama
146140  Rodolpho Teófilo - O Legado de um Pioneiro
Documentary
146141          Dankyavar Danka

```

```

Comedy
146143          Chico Albuquerque - Revelações
Documentary

[140736 rows x 2 columns]

# confirming if our dataset contains any missing value
idmb_unzip1.isna().sum()

primary_title    0
genres           0
dtype: int64

# Lets check for duplicates
idmb_unzip1.duplicated().sum()

1612

# the dataset contains 1612 duplicates we drop them keeping first
cleaned_imdb = idmb_unzip1.drop_duplicates()
cleaned_imdb

          primary_title
genres
0                      Sunghursh
Action, Crime, Drama
1      One Day Before the Rainy Season
Biography, Drama
2      The Other Side of the Wind
Drama
3      Sabse Bada Sukh
Comedy, Drama
4      The Wandering Soap Opera
Comedy, Drama, Fantasy
...
...
146138      The Secret of China
Adventure, History, War
146139      Kuambil Lagi Hatiku
Drama
146140      Rodolpho Teóphilo - O Legado de um Pioneiro
Documentary
146141      Dankyavar Danka
Comedy
146143      Chico Albuquerque - Revelações
Documentary

[139124 rows x 2 columns]

# confirming duplicates
cleaned_imdb.duplicated().sum()

```

0

Our imdb dataset is clean with a new name as cleaned_imdb now we work for our csv data that bom.movie_gross.csv. we start by reading it

```
bom_movie = pd.read_csv("zippedData/bom.movie_gross.csv")
bom_movie.head()
```

	title	studio	domestic_gross
0	Toy Story 3	BV	415000000.0
1	Alice in Wonderland (2010)	BV	334200000.0
2	Harry Potter and the Deathly Hallows Part 1	WB	296000000.0
3	Inception	WB	292600000.0
4	Shrek Forever After	P/DW	238700000.0

	foreign_gross	year
0	652000000	2010
1	691300000	2010
2	664300000	2010
3	535700000	2010
4	513900000	2010

```
# checking how our dataset looks like
bom_movie.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3387 entries, 0 to 3386
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   title                  3387 non-null   object  
1   studio                 3382 non-null   object  
2   domestic_gross         3359 non-null   float64  
3   foreign_gross          2037 non-null   object  
4   year                   3387 non-null   int64  
dtypes: float64(1), int64(1), object(3)
memory usage: 132.4+ KB
```

Data analysis

1. First we will convert the foreign_gross from object dtype to float dtype
2. Check the null or missing values and fill them, and drop where need be
3. Check and drop duplicates
4. Check and drop outliers

```
#checking the colums
```

```
bom_movie.columns
```

```
Index(['title', 'studio', 'domestic_gross', 'foreign_gross', 'year'],  
      dtype='object')
```

```
# we need the foreign gross column in as a float dtype. we need to  
convert it from object to float dtype
```

```
bom_movie['foreign_gross'] =
```

```
bom_movie['foreign_gross'].str.replace(',', '').astype(float)
```

```
bom_movie.describe()
```

	domestic_gross	foreign_gross	year
count	3.359000e+03	2.037000e+03	3387.000000
mean	2.874585e+07	7.487281e+07	2013.958075
std	6.698250e+07	1.374106e+08	2.478141
min	1.000000e+02	6.000000e+02	2010.000000
25%	1.200000e+05	3.700000e+06	2012.000000
50%	1.400000e+06	1.870000e+07	2014.000000
75%	2.790000e+07	7.490000e+07	2016.000000
max	9.367000e+08	9.605000e+08	2018.000000

```
# checking for missing values
```

```
bom_movie.isna().sum()
```

```
title          0  
studio         5  
domestic_gross 28  
foreign_gross 1350  
year           0  
dtype: int64
```

```
# The number of missing values in foreign _gross is a huge percentage  
to be dropped
```

```
# We calculate the mean value and use it to fill the missing values
```

```
mean_value = bom_movie['foreign_gross'].mean()
```

```
mean_value
```

```
74872810.15046637
```

```
# Filling the missing values using the mean
```

```
bom_movie['foreign_gross'] =
```

```
bom_movie['foreign_gross'].fillna(mean_value)
```

```
# We calculate the mean value and use it to fill the missing values
```

```
mean_value = bom_movie['domestic_gross'].mean()
```

```
mean_value
```

```
28745845.06698422
```

```
#Filling the missing values using the mean
bom_movie['domestic_gross'] =
bom_movie['domestic_gross'].fillna(mean_value)
```

```
# dropping miss values in studio
bom_movie1 = bom_movie.dropna(subset=['studio'])
bom_movie1
```

	title	studio
domestic_gross \		
0	Toy Story 3	BV
415000000.0		
1	Alice in Wonderland (2010)	BV
334200000.0		
2	Harry Potter and the Deathly Hallows Part 1	WB
296000000.0		
3	Inception	WB
292600000.0		
4	Shrek Forever After	P/DW
238700000.0		
...
...		
3382	The Quake	Magn.
6200.0		
3383	Edward II (2018 re-release)	FM
4800.0		
3384	El Pacto	Sony
2500.0		
3385	The Swan	Synergetic
2400.0		
3386	An Actor Prepares	Grav.
1700.0		

	foreign_gross	year
0	6.520000e+08	2010
1	6.913000e+08	2010
2	6.643000e+08	2010
3	5.357000e+08	2010
4	5.139000e+08	2010
...
3382	7.487281e+07	2018
3383	7.487281e+07	2018
3384	7.487281e+07	2018
3385	7.487281e+07	2018
3386	7.487281e+07	2018

```
[3382 rows x 5 columns]
```

```
bom_movie1.isna().sum()
```

```

title          0
studio         0
domestic_gross 0
foreign_gross  0
year           0
dtype: int64


# Checking for duplicates, bom_mvie1 has no duplicates
bom_mvie1.duplicated().sum()

0

# Checking for outliers domestic_gross column using the boxplot
sns.boxplot(x = bom_mvie1['domestic_gross'])

<Axes: xlabel='domestic_gross'>

```



```

# i will use the interquatile range to calculate and filter out
outliers
# Calculate the IQR
Q1 = bom_mvie1['domestic_gross'].quantile(0.25)
Q3 = bom_mvie1['domestic_gross'].quantile(0.75)
IQR = Q3 - Q1

```



```
# Define the bounds for outliers
```

```
lower_bound = Q1 - 1.5 * IQR
```

```
upper_bound = Q3 + 1.5 * IQR
```

```
# Filter out outliers
```

```
fi_bom_movie = bom_movie1[(bom_movie1['domestic_gross'] >=
lower_bound) & (bom_movie1['domestic_gross'] <= upper_bound)]
fi_bom_movie
```


	title	studio	domestic_gross
foreign_gross \			
21	Resident Evil: Afterlife	SGem	60100000.0
2.401000e+08			
25	The Tourist	Sony	67600000.0
2.107000e+08			
30	Gulliver's Travels	Fox	42800000.0
1.946000e+08			
34	The Sorcerer's Apprentice	BV	63200000.0
1.521000e+08			
44	Step Up 3-D	BV	42400000.0
1.169000e+08			
...
...			
3382	The Quake	Magn.	6200.0
7.487281e+07			
3383	Edward II (2018 re-release)	FM	4800.0
7.487281e+07			
3384	El Pacto	Sony	2500.0
7.487281e+07			
3385	The Swan	Synergetic	2400.0
7.487281e+07			
3386	An Actor Prepares	Grav.	1700.0
7.487281e+07			

	year
21	2010
25	2010
30	2010
34	2010
44	2010
...	...
3382	2018
3383	2018
3384	2018
3385	2018
3386	2018

```
[2984 rows x 5 columns]
```


```
# Checking outliers for the domestic_gross column after filtering out,  
a way of confirming  
sns.boxplot(x = fi_bom_movie ['domestic_gross'])
```

```
<Axes: xlabel='domestic_gross'>
```



```
# Checking out outliers in the foreign_gross column using the boxplot  
sns.boxplot(x=fi_bom_movie['foreign_gross'])
```

```
<Axes: xlabel='foreign_gross'>
```



```
# i will use the interquatile range to calculate and filter out
outliers
# Calculate the IQR
Q1 = fi_bom_movie['foreign_gross'].quantile(0.25)
Q3 = fi_bom_movie['foreign_gross'].quantile(0.75)
IQR = Q3 - Q1

# Define the bounds for outliers
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR

# Filter out outliers
cleaned_bom_movie = fi_bom_movie[(fi_bom_movie['foreign_gross'] >=
lower_bound) & (fi_bom_movie['foreign_gross'] <= upper_bound)]
cleaned_bom_movie
```

	title	studio \
34	The Sorcerer's Apprentice	BV
44	Step Up 3-D	BV
48	Legend of the Guardians: The Owls of Ga'Hoole	WB
49	The Wolfman	Uni.
50	The Bounty Hunter	Sony
...
3382	The Quake	Magn.

3383	Edward II (2018 re-release)	FM
3384	El Pacto	Sony
3385	The Swan	Synergetic
3386	An Actor Prepares	Grav.


	domestic_gross	foreign_gross	year
34	63200000.0	1.521000e+08	2010
44	42400000.0	1.169000e+08	2010
48	55700000.0	8.440000e+07	2010
49	62000000.0	7.780000e+07	2010
50	67099999.0	6.930000e+07	2010
...
3382	6200.0	7.487281e+07	2018
3383	4800.0	7.487281e+07	2018
3384	2500.0	7.487281e+07	2018
3385	2400.0	7.487281e+07	2018
3386	1700.0	7.487281e+07	2018

[2936 rows x 5 columns]

Checking outliers for the foreign_gross column after filtering out, a way of confirming

`sns.boxplot(x=cleaned_bom_movie['foreign_gross'])`

<Axes: xlabel='foreign_gross'>



```
df = pd.read_csv("zippedData/tn.movie_budgets.csv")
df.head()
```

	id	release_date	movie \
0	1	Dec 18, 2009	Avatar
1	2	May 20, 2011	Pirates of the Caribbean: On Stranger Tides
2	3	Jun 7, 2019	Dark Phoenix
3	4	May 1, 2015	Avengers: Age of Ultron
4	5	Dec 15, 2017	Star Wars Ep. VIII: The Last Jedi

	production_budget	domestic_gross	worldwide_gross
0	\$425,000,000	\$760,507,625	\$2,776,345,279
1	\$410,600,000	\$241,063,875	\$1,045,663,875
2	\$350,000,000	\$42,762,350	\$149,762,350
3	\$330,600,000	\$459,005,868	\$1,403,013,963
4	\$317,000,000	\$620,181,382	\$1,316,721,747

#checking the shape of the datasets it 5782 contains rows and 6 columns

```
df.shape
```

```
(5782, 6)
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 5782 entries, 0 to 5781
```

```
Data columns (total 6 columns):
```

#	Column	Non-Null Count	Dtype
0	id	5782 non-null	int64
1	release_date	5782 non-null	object
2	movie	5782 non-null	object
3	production_budget	5782 non-null	object
4	domestic_gross	5782 non-null	object
5	worldwide_gross	5782 non-null	object

```
dtypes: int64(1), object(5)
```

```
memory usage: 271.2+ KB
```

Data analysis

1. First we will convert the production_budget, domestic_gross, worldwide_gross from object dtype to float dtype
2. Check the null or missing values and fill them, and drop where need be
3. Check and drop duplicates
4. Check and drop outliers

To convert the production_budget, domestic_gross, worldwide_gross we have to get rid of commas and the dollar signs. The reason been to work with numbers in calculations or data analysis, you need to get rid of these symbols since computers will not treat them as numbers.

```
# we defined a function
def remove_dollar_sign(x):
    return x.replace('$', '').replace(',', '')

# defined the columns we need to remove symbols and commas
remove_columns = ['production_budget', 'domestic_gross',
'worldwide_gross']
# we use for loops to iterate through the columns
for column in remove_columns:
    df[column] = df[column].apply(remove_dollar_sign)

# confirming if our columns contains the signs, good we are okay
df.head(2)
```

	id	release_date	movie
0	1	Dec 18, 2009	Avatar
1	2	May 20, 2011	Pirates of the Caribbean: On Stranger Tides

	production_budget	domestic_gross	worldwide_gross
0	425000000	760507625	2776345279
1	410600000	241063875	1045663875

```
# first we convert the production_budget, domestic_gross,
worldwide_gross from object dtype to float dtype
# converting the production_budget
df['production_budget'] = df['production_budget'].str.replace(',',
 '').astype(float)
```

```
# converting the domestic_gross
df['domestic_gross'] = df['domestic_gross'].str.replace(',',
 '').astype(float)
```

```
#converting the worldwide_gross
df['worldwide_gross'] = df['worldwide_gross'].str.replace(',',
 '').astype(float)
```

```
#successfully converted to float dtype
df.dtypes
```

id	int64
release_date	object
movie	object
production_budget	float64
domestic_gross	float64
worldwide_gross	float64
dtype:	object

```
# checking the stats i.e mean, median, mode, count e.t.c
df.describe()
```

	id	production_budget	domestic_gross	worldwide_gross
count	5782.000000	5.782000e+03	5.782000e+03	5.782000e+03

mean	50.372363	3.158776e+07	4.187333e+07	9.148746e+07
std	28.821076	4.181208e+07	6.824060e+07	1.747200e+08
min	1.000000	1.100000e+03	0.000000e+00	0.000000e+00
25%	25.000000	5.000000e+06	1.429534e+06	4.125415e+06
50%	50.000000	1.700000e+07	1.722594e+07	2.798445e+07
75%	75.000000	4.000000e+07	5.234866e+07	9.764584e+07
max	100.000000	4.250000e+08	9.366622e+08	2.776345e+09

performing sanity check

```
#checking missing values, drop, or fill them.
df.isna().sum() # our data has no missing values

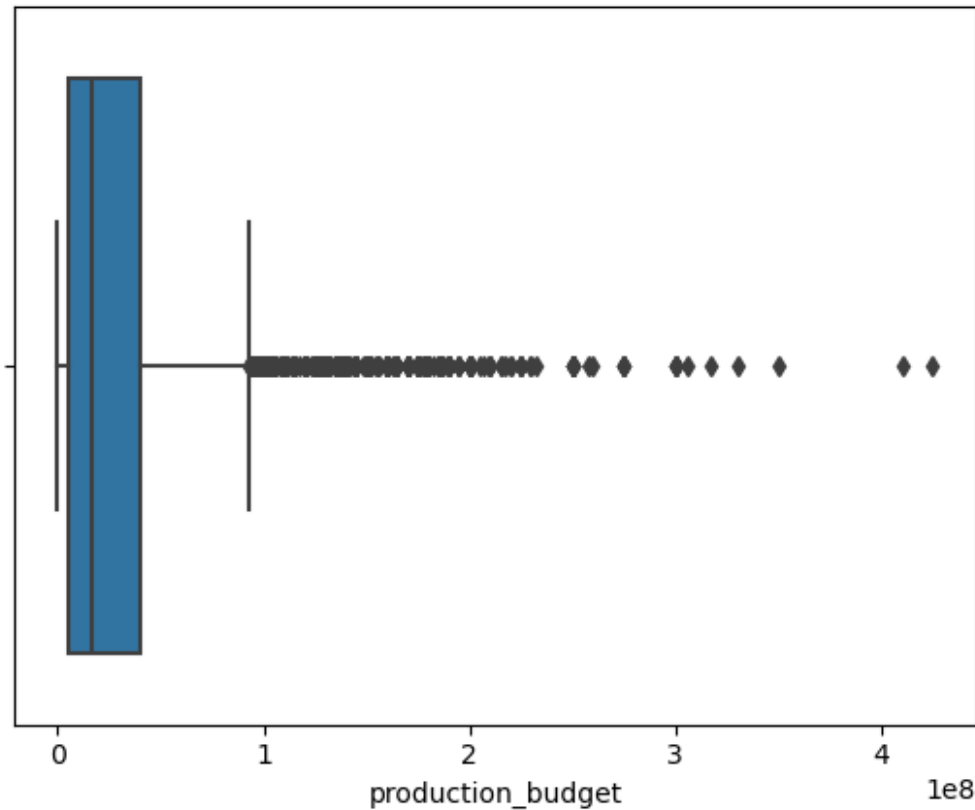
id                0
release_date      0
movie             0
production_budget 0
domestic_gross    0
worldwide_gross   0
dtype: int64

# checking for duplicates
df.duplicated().sum() # no duplicates

0

# checking for outliers using sns library and box plot
sns.boxplot(x =df['production_budget'])

<Axes: xlabel='production_budget'>
```



```
# i will use the interquatile range to calculate and filter out outliers
```

```
# Calculate the IQR
```

```
Q1 = df['production_budget'].quantile(0.25)
```

```
Q3 = df['production_budget'].quantile(0.75)
```

```
IQR = Q3 - Q1
```

```
# Define the bounds for outliers
```

```
lower_bound = Q1 - 1.5 * IQR
```

```
upper_bound = Q3 + 1.5 * IQR
```

```
# Filter out outliers
```

```
df1 = df[(df['production_budget'] >= lower_bound) &
```

```
(df['production_budget'] <= upper_bound)]
```

```
df1
```

	id	release_date	movie
production_budget \			
431	32	Feb 14, 2008	The Spiderwick Chronicles
92500000.0			
432	33	Nov 5, 2004	The Incredibles
92000000.0			
433	34	Feb 14, 2013	A Good Day to Die Hard
92000000.0			


434	35	Apr 9, 2004	The Alamo
92000000.0			
435	36	Dec 22, 1995	Cutthroat Island
92000000.0			
...
..			
5777	78	Dec 31, 2018	Red 11
7000.0			
5778	79	Apr 2, 1999	Following
6000.0			
5779	80	Jul 13, 2005	Return to the Land of Wonders
5000.0			
5780	81	Sep 29, 2015	A Plague So Pleasant
1400.0			
5781	82	Aug 5, 2005	My Date With Drew
1100.0			

	domestic_gross	worldwide_gross
431	71195053.0	162839667.0
432	261441092.0	614726752.0
433	67349198.0	304249198.0
434	22406362.0	23911362.0
435	10017322.0	18517322.0
...
5777	0.0	0.0
5778	48482.0	240495.0
5779	1338.0	1338.0
5780	0.0	0.0
5781	181041.0	181041.0


[5351 rows x 6 columns]

```
sns.boxplot(x =df1['production_budget'])
```

```
<Axes: xlabel='production_budget'>
```



```
sns.boxplot(x =df1['domestic_gross'])  
<Axes: xlabel='domestic_gross'>
```



```
# i will use the interquatile range to calculate and filter out outliers
```

```
# Calculate the IQR
```

```
Q1 = df1['domestic_gross'].quantile(0.25)
```

```
Q3 = df1['domestic_gross'].quantile(0.75)
```

```
IQR = Q3 - Q1
```

```
# Define the bounds for outliers
```

```
lower_bound = Q1 - 1.5 * IQR
```

```
upper_bound = Q3 + 1.5 * IQR
```

```
# Filter out outliers
```

```
df2 = df1[(df1['domestic_gross'] >= lower_bound) &
```

```
(df1['domestic_gross'] <= upper_bound)]
```

```
df2
```

	id	release_date	movie
production_budget \			
431	32	Feb 14, 2008	The Spiderwick Chronicles
92500000.0			
433	34	Feb 14, 2013	A Good Day to Die Hard
92000000.0			
434	35	Apr 9, 2004	The Alamo
92000000.0			

```

435    36  Dec 22, 1995          Cutthroat Island
92000000.0
436    37  Dec 25, 2013  The Secret Life of Walter Mitty
91000000.0
...    ..          ...
...
5777   78  Dec 31, 2018          Red 11
7000.0
5778   79   Apr 2, 1999          Following
6000.0
5779   80  Jul 13, 2005  Return to the Land of Wonders
5000.0
5780   81  Sep 29, 2015          A Plague So Pleasant
1400.0
5781   82   Aug 5, 2005          My Date With Drew
1100.0


```

	domestic_gross	worldwide_gross
431	71195053.0	162839667.0
433	67349198.0	304249198.0
434	22406362.0	23911362.0
435	10017322.0	18517322.0
436	58236838.0	187861183.0
...
5777	0.0	0.0
5778	48482.0	240495.0
5779	1338.0	1338.0
5780	0.0	0.0
5781	181041.0	181041.0


```
[4994 rows x 6 columns]
```

```
sns.boxplot(x =df2['domestic_gross'])
```

```
<Axes: xlabel='domestic_gross'>
```



```
sns.boxplot(x =df2['worldwide_gross'])  
<Axes: xlabel='worldwide_gross'>
```



```
# i will use the interquatile range to calculate and filter out outliers
```

```
# Calculate the IQR
```

```
Q1 = df2['worldwide_gross'].quantile(0.25)
```

```
Q3 = df2['worldwide_gross'].quantile(0.75)
```

```
IQR = Q3 - Q1
```

```
# Define the bounds for outliers
```

```
lower_bound = Q1 - 1.5 * IQR
```

```
upper_bound = Q3 + 1.5 * IQR
```

```
# Filter out outliers
```

```
tn_movie = df2[(df2['worldwide_gross'] >= lower_bound) &  
(df2['worldwide_gross'] <= upper_bound)]
```

```
tn_movie
```

	id	release_date	movie
production_budget \			
434	35	Apr 9, 2004	The Alamo
92000000.0			
435	36	Dec 22, 1995	Cutthroat Island
92000000.0			
466	67	Nov 21, 2001	Spy Game
90000000.0			


468	69	Mar 10, 2000	Mission to Mars
900000000.0			
470	71	Dec 17, 1999	Bicentennial Man
900000000.0			
...
..			
5777	78	Dec 31, 2018	Red 11
7000.0			
5778	79	Apr 2, 1999	Following
6000.0			
5779	80	Jul 13, 2005	Return to the Land of Wonders
5000.0			
5780	81	Sep 29, 2015	A Plague So Pleasant
1400.0			
5781	82	Aug 5, 2005	My Date With Drew
1100.0			

	domestic_gross	worldwide_gross
434	22406362.0	23911362.0
435	10017322.0	18517322.0
466	62362560.0	143049560.0
468	60874615.0	106000000.0
470	58220776.0	87420776.0
...
5777	0.0	0.0
5778	48482.0	240495.0
5779	1338.0	1338.0
5780	0.0	0.0
5781	181041.0	181041.0

[4657 rows x 6 columns]

```
sns.boxplot(x =tn_movie['worldwide_gross'])
```

```
<Axes: xlabel='worldwide_gross'>
```



Our dataset is clean and the new name is cleaned_bom_movie

By merging the datasets I will create a consolidated dataset for a better and not complicated analysis. To ensure data integrity and completeness, I will exclude any records that do not have a match between the datasets, thus eliminating any potential missing values for the features.

1. I will create a unique key that is a primary key. in the two datasets they have a common column that contains the title of the movies just that they have different column names. i will rename the column names in both dataset
2. Having a primary key named movie_title i can proceed on merging the two datasets

```
# rename primary title to more readable and understanding column name
cleaned_imdb = cleaned_imdb.rename(columns={'primary_title':
'movie_title'})
cleaned_imdb
```

	movie_title
genres	
0	Sunghursh
Action, Crime, Drama	
1	One Day Before the Rainy Season


```

Biography,Drama
2                                The Other Side of the Wind
Drama
3                                Sabse Bada Sukh
Comedy,Drama
4                                The Wandering Soap Opera
Comedy,Drama,Fantasy
...
...
146138                           The Secret of China
Adventure,History,War
146139                           Kuambil Lagi Hatiku
Drama
146140 Rodolpho Teóphilo - O Legado de um Pioneiro
Documentary
146141                           Dankyavar Danka
Comedy
146143                           Chico Albuquerque - Revelações
Documentary

```

```
[139124 rows x 2 columns]
```

```
cleaned_imdb.columns
```

```
Index(['movie_title', 'genres'], dtype='object')
```

```
# renaming the title in the bom dataset to movie
```

```
cleaned_bom_movie =
```

```
cleaned_bom_movie.rename(columns={'title':'movie_title'})
```

```
cleaned_bom_movie
```

	movie_title	studio \
34	The Sorcerer's Apprentice	BV
44	Step Up 3-D	BV
48	Legend of the Guardians: The Owls of Ga'Hoole	WB
49	The Wolfman	Uni.
50	The Bounty Hunter	Sony
...
3382	The Quake	Magn.
3383	Edward II (2018 re-release)	FM
3384	El Pacto	Sony
3385	The Swan	Synergetic
3386	An Actor Prepares	Grav.

	domestic_gross	foreign_gross	year
34	63200000.0	1.521000e+08	2010
44	42400000.0	1.169000e+08	2010
48	55700000.0	8.440000e+07	2010
49	62000000.0	7.780000e+07	2010
50	67099999.0	6.930000e+07	2010

```

...
3382      6200.0    7.487281e+07    2018
3383      4800.0    7.487281e+07    2018
3384      2500.0    7.487281e+07    2018
3385      2400.0    7.487281e+07    2018
3386      1700.0    7.487281e+07    2018

```

```
[2936 rows x 5 columns]
```

```
cleaned_bom_movie.columns
```

```
Index(['movie_title', 'studio', 'domestic_gross', 'foreign_gross',
      'year'], dtype='object')
```

```
cleaned_tn_movie = tn_movie.rename(columns={'movie': 'movie_title'})
cleaned_tn_movie
```

```

      id  release_date      movie_title
production_budget \
434   35   Apr 9, 2004      The Alamo
92000000.0
435   36  Dec 22, 1995      Cutthroat Island
92000000.0
466   67  Nov 21, 2001      Spy Game
90000000.0
468   69  Mar 10, 2000      Mission to Mars
90000000.0
470   71  Dec 17, 1999      Bicentennial Man
90000000.0
...    ..          ...
..
5777  78  Dec 31, 2018      Red 11
7000.0
5778  79   Apr 2, 1999      Following
6000.0
5779  80  Jul 13, 2005  Return to the Land of Wonders
5000.0
5780  81  Sep 29, 2015      A Plague So Pleasant
1400.0
5781  82   Aug 5, 2005      My Date With Drew
1100.0

```

```

      domestic_gross  worldwide_gross
434      22406362.0      23911362.0
435      10017322.0      18517322.0
466      62362560.0      143049560.0
468      60874615.0      106000000.0
470      58220776.0      87420776.0
...
5777           0.0           0.0

```

5778	48482.0	240495.0
5779	1338.0	1338.0
5780	0.0	0.0
5781	181041.0	181041.0

[4657 rows x 6 columns]

We will use the merge method to merge cleaned_idmb and cleaned_bom_movie cleaned_tn_movie on movie_title as our primary key using an inner join parameter and we name the merged dataset as movies

```
# Merge tn_mb and bom_movie_gross on movie_title
```

```
movies1 = pd.merge(cleaned_idmb,
                    cleaned_bom_movie,
                    on=['movie_title'],
                    how='inner')
```

```
# Checking the 5 first rows as well as confirming if our dataset has merged
```

```
movies1.head()
```

	movie_title	genres	studio \
0	Wazir	Action, Crime, Drama	Relbig.
1	On the Road	Adventure, Drama, Romance	IFC
2	On the Road	Drama	IFC
3	On the Road	Documentary	IFC
4	The Secret Life of Walter Mitty	Adventure, Comedy, Drama	Fox

	domestic_gross	foreign_gross	year
0	1100000.0	7.487281e+07	2016
1	744000.0	8.000000e+06	2012
2	744000.0	8.000000e+06	2012
3	744000.0	8.000000e+06	2012
4	58200000.0	1.299000e+08	2013

```
# Merge movies1 and tn.movie on movie_title
```

```
movies2 = pd.merge(movies1,
                    cleaned_tn_movie,
                    on=['movie_title'],
                    how='inner')
```

```
movies2.head()
```

	movie_title	genres	studio \
0	On the Road	Adventure, Drama, Romance	IFC

1	On the Road	Drama	IFC
2	On the Road	Documentary	IFC
3	A Walk Among the Tombstones	Action, Crime, Drama	Uni.
4	The Rum Diary	Comedy, Drama	FD

	domestic_gross_x	foreign_gross	year	id	release_date
production_budget \					
0	744000.0	8000000.0	2012	17	Mar 22, 2013
25000000.0					
1	744000.0	8000000.0	2012	17	Mar 22, 2013
25000000.0					
2	744000.0	8000000.0	2012	17	Mar 22, 2013
25000000.0					
3	26300000.0	26900000.0	2014	67	Sep 19, 2014
28000000.0					
4	13100000.0	10800000.0	2011	16	Oct 28, 2011
45000000.0					

	domestic_gross_y	worldwide_gross
0	720828.0	9313302.0
1	720828.0	9313302.0
2	720828.0	9313302.0
3	26017685.0	62108587.0
4	13109815.0	21544732.0

Our dataset has successfully merged but it contains some columns that we will not use so let's drop them they include: domestic_gross_x, foreign_gross, year, id, release_date

```
movies = movies2.drop(['domestic_gross_x', 'foreign_gross', 'id',
                        'release_date'], axis=1)
movies.head(2)
```

	movie_title	genres	studio	year
production_budget \				
0	On the Road	Adventure, Drama, Romance	IFC	2012
25000000.0				
1	On the Road	Drama	IFC	2012
25000000.0				

	domestic_gross_y	worldwide_gross
0	720828.0	9313302.0
1	720828.0	9313302.0

```
movies.columns
```

```
Index(['movie_title', 'genres', 'studio', 'year', 'production_budget',
       'domestic_gross_y', 'worldwide_gross'],
      dtype='object')
```

```
movies.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1037 entries, 0 to 1036
Data columns (total 7 columns):
#   Column                Non-Null Count  Dtype
---  -
0   movie_title            1037 non-null   object
1   genres                 1037 non-null   object
2   studio                 1037 non-null   object
3   year                   1037 non-null   int64
4   production_budget      1037 non-null   float64
5   domestic_gross_y       1037 non-null   float64
6   worldwide_gross        1037 non-null   float64
dtypes: float64(3), int64(1), object(3)
memory usage: 56.8+ KB

movies.shape

(1037, 7)

```

Our final dataset is called movies, it is a clean data, note we cleaned the datasets before merging them. Our dataset contains 1037 rows and 6 columns

Data visualization

The aim of starting a business is to make profit, Voice of Kenya want to venture in the movie production business. The question is, is this movie production profitable. Our dataset has two columns i.e domestic_gross_y and worldwide_gross where worldwide_gross is total revenue collected worldwide that means that even the domestic_gross is inclusive

profitability calculation or return on investment

```

# Removing a film that is multiplred genred like the same movie has
# multiple genre i.e adventure, drama, comedy
movies['genres'] = movies['genres'].str.split(',').str[0].str.strip()
len(movies)

```

```
1037
```

```

# we have 67 duplicates which need dropping
movies.duplicated().sum()

```

```
67
```

```

# drop duplicates this duplicates arised when we splitted our dataset
# genre which had multiple genre and we kept one
movies_1 = movies.drop_duplicates()
movies_1.duplicated().sum()

```

```
0
```

lets calculate the profit distribution per genre and draw a visualization to see which genre is more profitable

```
# calculating profit
```

```
movies_1['profits'] = movies_1['worldwide_gross'] -  
movies_1['production_budget']
```

```
C:\Users\PC\AppData\Local\Temp\ipykernel_3088\3027505178.py:2:
```

```
SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame.
```

```
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation:
```

```
https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#  
returning-a-view-versus-a-copy
```

```
movies_1['profits'] = movies_1['worldwide_gross'] -  
movies_1['production_budget']
```

```
# visualizing the distribution of profits
```

```
# using bargraph
```

```
plt.figure(figsize=(14,8))
```

```
sns.barplot(x = 'genres', y = 'profits', data = movies_1 , ci= None)
```

```
plt.title('Profit distribution per genre')
```

```
plt.xlabel('genres')
```

```
plt.xticks(rotation=45 )
```

```
plt.ylabel('profits')
```


```
C:\Users\PC\AppData\Local\Temp\ipykernel_3088\3856034616.py:4:
```

```
FutureWarning:
```

```
The `ci` parameter is deprecated. Use `errorbar=None` for the same  
effect.
```

```
sns.barplot(x = 'genres', y = 'profits', data = movies_1 , ci= None)
```

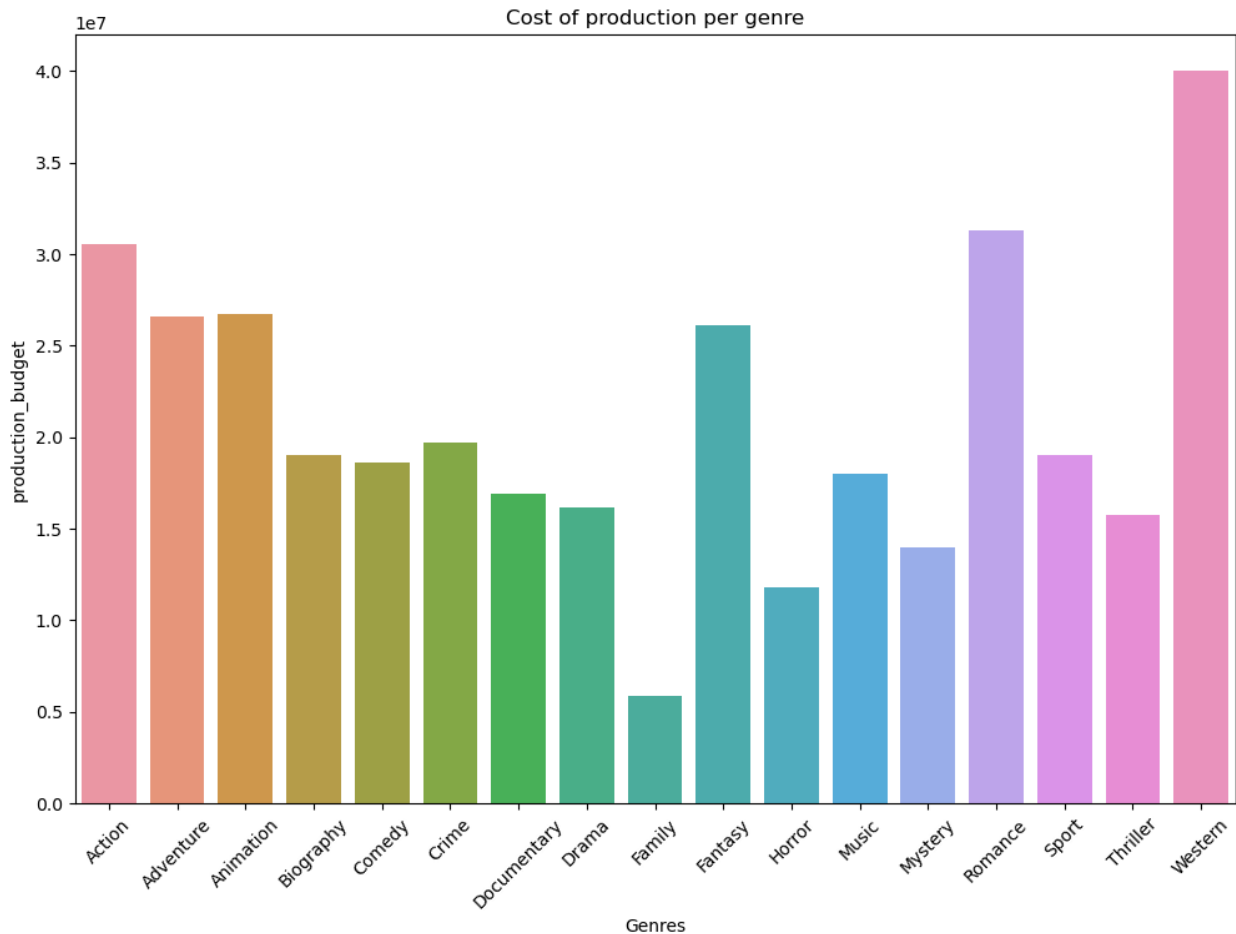
```
Text(0, 0.5, 'profits')
```



Lets do the production of movie budget and the genre. in production the cost differs with the genre for example animation is cheaper compared to adventure.

```
# Calculate the average production per genre
movies_2 = movies_1.groupby('genres')
['production_budget'].mean().reset_index()

# plotting bargraph
plt.figure(figsize=(12,8))
sns.barplot(x = 'genres', y = 'production_budget', data = movies_2)
plt.title('Cost of production per genre')
plt.xlabel('Genres')
plt.xticks(rotation=45)
plt.ylabel('production_budget')
plt.show()
```



Lets do the production of movie budget and the genre. in production the cost differs with the genre for example animation is cheaper compared to adventure.

Define X and Y

importing linear regression libraries, we want to see the relationship between the domestic_gross and worldwide gross if they have any relationship

```
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

# performing a simple linear regression
X = movies_1['domestic_gross_y']
y = movies_1['worldwide_gross']
# Add a constant to the independent variable
X = sm.add_constant(X)
# Perform linear regression
model = sm.OLS(y, X).fit()
#summary of the regression
print(model.summary())
```


OLS Regression Results

```

=====
Dep. Variable:      worldwide_gross    R-squared:
0.731
Model:              OLS                Adj. R-squared:
0.731
Method:             Least Squares      F-statistic:
2630.
Date:               Sun, 28 Jul 2024   Prob (F-statistic):
3.11e-278
Time:               10:44:49          Log-Likelihood:
-17621.
No. Observations:   970              AIC:
3.525e+04
Df Residuals:       968              BIC:
3.526e+04
Df Model:           1
Covariance Type:    nonrobust

```

```

=====
=====
               coef      std err          t      P>|t|
[0.025      0.975]
-----
const          7.445e+06    8.73e+05     8.526     0.000
5.73e+06    9.16e+06
domestic_gross_y  1.5875      0.031    51.286     0.000
1.527      1.648
=====
=====

```

```

=====
Omnibus:          337.901    Durbin-Watson:
1.624
Prob(Omnibus):    0.000    Jarque-Bera (JB):
1239.045
Skew:             1.656    Prob(JB):
8.81e-270
Kurtosis:         7.437    Cond. No.
4.09e+07
=====
=====


```

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 4.09e+07. This might indicate that

there are strong multicollinearity or other numerical problems.

```
# Scatter plot to visualize the relationship and a regression plot
plt.figure(figsize=(10, 6))
sns.scatterplot(x='domestic_gross_y', y='worldwide_gross',
data=movies_1)
sns.regplot(x='domestic_gross_y', y='worldwide_gross', data=movies_1,
line_kws={'color': 'red'})
plt.title('Domestic Gross vs Worldwide Gross')
plt.xlabel('Domestic Gross')
plt.ylabel('Worldwide Gross')
plt.show()
```



Recommendations

1. Profitable Genre - Voice of Kenya should focus more on production of mystery genre movie. This genre has earned about 100m in a period of 8 years. Voice of Kenya should also venture in producing adventure, action, comedy, horror, and thriller films. These genres offer a balance of manageable profitability, making them the most strategic investment for achieving strong and wide market.
2. Cost of production - This is a key aspect. For example producing the western genre is very costly yet it is not profitable at all. In our data it shows that the cost of production of the thriller is not very high but it gives the best profit Based on the

analysis of production costs per genre. Producing Horror, Comedy, and Thriller films, which offer a good balance of moderate production costs and high profitability.

3. Using the linear regression model, it shows that there is a positive linear relationship between domestic and worldwide gross. The strong positive linear relationship between Domestic Gross and Worldwide Gross suggests that movies performing well domestically tend to also perform well internationally. Invest in genres and film types that have historically performed well domestically. Focus on marketing strategies to boost domestic performance, as it is a strong predictor of worldwide success

Conclusion

Voice of Kenya should focus on producing Mystery, Adventure, Action, Comedy, Horror, and Thriller films for optimal profitability and market reach. Emphasize cost-effective genres like Thriller, Horror, and Comedy. Leverage the strong positive relationship between domestic and worldwide gross by enhancing domestic performance through strategic marketing. Avoid high-cost, low-profit genres like Western.